

Full Stack ML/AI Training Track

By Dr.Vishwanath Rao

LEVEL 1

Pre-requisites

- Fundamentals of any programming language
- Familiarity with basic Algebra, Statistic, Probability and Calculus (Recommended)

Track numbers

- Track 1, 2 and 3

Learning Outcome

After going through these tracks, learners will:

- have built a solid foundation for any ML/AI development. - be familiar with most of the prominent libraries and frameworks for ML implementations.

Capstone Project

- 1) Binary Search - Create an application that will try to guess a number a user chooses between from a predefined set.
 - 2) Image Resizer - build an application that can resize an image file by changing its dimensions to the target parameters
- After going through these tracks, learners will:

Level 2

Pre-requisites

Must have completed L2 or its equivalent

Track numbers

- Track 4 and 5

Learning Outcome

After going through these tracks, learners will:

- gain an in-depth understanding of most of learning models and their use cases; hyperparameter tuning and ensemble techniques to optimize the learning models or algorithms. - be able to use the prominent third-party ML services provided by various cloud vendors and implement real (industry) use cases.

Capstone Project

1) Predict Housing Price - Predict the selling price of a new home in a given region

2) Chatbots - Create a chatbot for customer service.

After going through these tracks, learners will:

- learn specialized AI techniques and applications. - be able to understand and contribute to the R&D of the specified fields

Level 3

Pre-requisites

Must have completed L1 or its equivalent

Track numbers

- Track 6 and 7

Learning Outcome

After going through these tracks, learners will:

- learn specialized AI techniques and applications. - be able to understand and contribute to the R&D of the specified fields.

Capstone Project

1) Facial Emotion Recognition and Detection - Use Deep Learning to identify and interpret human facial expressions.

2) Stock Price Predictor - This project seeks to utilize Deep Learning models, Long-Short Term Memory (LSTM) Neural Network algorithm, to predict stock prices.

Track 1: Math for AI/ML

Statistics

- Types of Data
- Levels of Measurement - Categorical Variables - Numerical Variables
- Histogram
- Cross Tables and

Scatter Plots

- Mean, Median, Mode - Skewness
- Variance
- Standard Deviation - Covariance
- Correlation

Linear Algebra

- What is a Matrix?
- Scalars and Vectors
- Linear Algebra and Geometry
- What is a Tensor?
- Addition and Subtraction of Matrices - Transpose of a Matrix
- Dot Product of Matrices
- Matrix Norms
- Special Matrices and Vectors

Eigenvalues and Eigenvectors - Principle component analysis - Singular value decomposition

Calculus

- Derivatives (Scalar Derivative-Chain rule), Partial and Directional Derivative.
- Integrals
- Gradients
- Differential Operators
- Convex Optimization
- Gradient algorithms- local/global maxima and minima, SGD, NAG, MAG, Adams

Probability

- Elements of Probability
- Random Variables
- Distributions (binomial, Bernoulli, Poisson, Exponential, Gaussian)
- Variance and Expectation - Bayes' Theorem, MAP, MLE - Special Random Variables - Markov Chain
- Information Theory

Track: 2 Fundamentals of Machine learning

Working with Data

- Exploring your data
- Cleaning and Munging
- Manipulating Data
- Rescaling
- Dimensionality Reduction

Machine Learning concepts

- What is Machine Learning? - Modelling
- Overfitting and Underfitting - The Bias-Variance tradeoff - Data Cleaning
- Data Preprocessing

Supervised Learning

- Simple Linear Regression - Multiple Regression
- Logistic Regression
- Logistic Regression
- Decision Tree

Unsupervised Learning

- Clustering
- K-means clustering
- KNN (k-nearest neighbors)

Track 3: Python for Implementing AI/ML

Python Basics

- Variables and Data Types - Basic python Syntax
- Operators
- Conditional Statements
- Functions - Sequence - Iterations - OOP
- Managing Data using python - Modules and Methods
- Essential Libraries
- Testing Python Code
- Debugging Python Applications

Databases concepts for implementing Machine Learning

Python Frameworks and libraries for ML

- Getting Started with Anaconda
- Installing and understanding the Jupyter's Interface - Prerequisites for Coding in the Jupyter Notebooks
- Numpy
- Scipy
- Scikit-learn
- Theano
- TensorFlow
- Keras
- PyTorch
- Pandas
- Matplotlib

Track 4: Advance Machine Learning

Advance Supervised Learning

- Simple Linear Regression - Multiple Regression
- Logistic Regression
- Logistic Regression
- Decision Tree

- Random Forest
- Naive Bayes
- Neural Networks - NLP
- Network Analysis

Unsupervised Learning

- Clustering
- K-means clustering
- KNN (k-nearest neighbors)
- Hierarchical clustering
- Anomaly detection
- Neural Networks
- Principle Component Analysis
- Independent Component Analysis - Apriori algorithm

Semi-supervised learning

Ensemble methods

- Bagging - Stacking - Boosting

Reinforcement Learning

- State-action-reward-state-action (SARSA) - Q-learning
- Deep Q-Networks

Deep Learning

Computational Learning Theory

Genetics Algorithm

The End-to-End Machine Learning Process (Putting it all together)

- Collection of Data
- Data Wrangling
- Model Selection
- Model Evaluation
- Model Deployment

Track 5: Machine Learning in the Cloud

Cloud Fundamentals

ML using ORACLE Cloud

Track 6: Artificial Intelligence Essentials

Introduction to Artificial Intelligence (AI)

- Elements Of Intelligence
- Intelligent Systems
- Research Areas of AI
- Agents and Environments
- Popular Search Algorithms
- Industry Applications of Artificial Intelligence - Fuzzy Logic Systems
- Expert systems

