Kubernetes Course Contents

By Dr. Vishwanath Rao

Kubernetes Core Concepts

- Kubernetes Basics
- What is Kubernetes?
- Container Orchestration
- Kubernetes Architecture
- Kubernetes Concepts
- Cluster and Namespace
- Nodes
- Master
- Pod
- Using Pods to Group Containers
- Label
- Label Syntax
- Annotation
- Label Selector
- Replication Controller and Replica Set
- Service
- Storage Volume
- Secret
- Resource Quota
- Authentication and Authorization
- Routing
- Docker Registry
- Summary

Kubernetes Architecture

- Architecture Diagram
- Components
- Kubernetes Cluster
- Master Node
- Kube-Control-Manager
- Nodes
- Other Components
- Interacting with Kubernetes

- Summary

Design

- Traditional Applications
- Virtual Machines
- Containerized Applications
- Decoupled Resources
- Transience
- Flexible Framework
- Application Resource Usage
- Measuring Resource Usage
- Docker Resource Usage Statistics
- Docker Container Resource Constraints
- Docker Run Command Resource Flags
- Using Label Selectors
- Equality Based Label Selector
- Set Based Label Selector
- Multi-Container Pods
- Sidecar Container
- Sidecar Container Uses
- Adapter Container
- Summary

Deployment Configuration

- Introduction to Volumes
- Container OS file system storage
- Docker Volumes
- Kubernetes Volumes
- Volume Specs
- K8S Volume Types
- Cloud Resource Types
- emptyDir
- Using an emptyDir Volume
- Other Volume Types
- Persistent Volumes
- Creating a Volume
- Persistent Volume Claim
- Persistent Volume
- Pod that uses Persistent Volume
- Dynamic Volume Provisioning
- Requesting Dynamic Storage
- Secrets
- Creating Secrets from Files
- Creating Secrets from Literals

- Using Secrets
- configMaps
- Creating configMaps from Literals
- Creating configMaps from files
- Using configMaps
- Security Context
- Security Context Usage
- Deployment Configuration Status
- Replicas
- Scaling
- Rolling Updates
- Summary

Managing Kubernetes Nodes

- Kubernetes Scheduler Overview
- Trusting the Kubernetes Scheduler
- Scheduling Process
- Scheduling Process - Predicates
- Scheduling Process - Priorities
- Scheduling Algorithm
- Kubernetes Scheduling Algorithm
- Scheduling Conflicts
- Controlling Scheduling
- Label Selectors
- Label Selectors (Contd.)
- Node Affinity and Anti-affinity
- Node Affinity Example
- Node Antiaffinity Example
- Taints and Tolerations
- Taints and Tolerations (Contd.)
- Taints and Tolerations - Example
- Summary

High-Availability K8S Clusters

- Overview
- Cluster Components
- Reliable Nodes
- Reliable Storage
- Replicated API Servers
- Load Balancing
- Cluster Control Components
- Installing Configuration Files
- Summary

Managing Kubernetes Clusters

- Kubernetes Cluster
- Points of Configuration
- Sizing a Cluster
- Resizing a Cluster
- Horizontal auto-scaling
- Maintenance on a Node
- Kubernetes Cluster Limits
- Considerations for Running Large Clusters
- Considerations for Running Multiple Clusters
- Multi-Cluster Architectural Choices
- Resource Quotas
- Example Resource Allocation
- Cluster Failure Modes
- Cluster Troubleshooting Workflow
- Summary

Exposing Applications

- Kubernetes Services
- Service Resources
- Service Type
- ClusterIP
- NodePort
- NodePort from Service Spec
- LoadBalancer
- LoadBalancer from Service Spec
- ExternalName
- Accessing Applications
- Service Without a Selector
- Ingress
- Ingress Resource Example
- Ingress Controller
- Service Mesh
- Summary

Troubleshooting Kubernetes

- Troubleshooting Overview
- Objects in Kubernetes
- Relationships in Kubernetes
- Operations in Kubernetes
- Understanding the Issue
- Troubleshooting Tools
- Troubleshooting Commands
- Troubleshooting Pods
- Troubleshooting the Cluster

- Cluster Failure Modes
- Monitoring
- Monitoring Applications
- Accessing Logs
- Logging Tools
- Conformance Testing
- Summary